



Goldmansachs Project

This project analyzes customer transaction behavior using the Goldman Sachs dataset.

The primary goal is to clean the data, understand customer activity patterns, and categorize accounts into Active, Inactive, and Dormant based on transaction frequency, recency, balance, and inflow trends.

Objectives:-

1. Clean and preprocess customer and transaction data.
2. Analyze account activity trends.
3. Define metrics that determine Active vs Inactive accounts.
4. Create new behavioral categories (frequency, inflow, balance).
5. Extract insights to help Goldman Sachs improve customer engagement.

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Data Loading and Initial Exploration

Here we load the dataset and inspect its structure. This step helps us understand:

Number of rows and columns

Data types

Missing values

Basic statistics of numerical variables

```
In [4]: df = pd.read_csv("goldman_sachs.csv")
df.head()
```

Out[4]:

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product
0	33	CUST6549	ACC12334	Credit	Withdrawal	Saving Account
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan

EDA

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionID          800 non-null    int64
1   CustomerID             800 non-null    object
2   AccountID              800 non-null    object
3   AccountType            800 non-null    object
4   TransactionType         800 non-null    object
5   Product                800 non-null    object
6   Firm                   800 non-null    object
7   Region                 800 non-null    object
8   Manager                800 non-null    object
9   TransactionDate         800 non-null    object
10  TransactionAmount       800 non-null    float64
11  AccountBalance          800 non-null    float64
12  RiskScore               800 non-null    float64
13  CreditRating            800 non-null    int64
14  TenureMonths            800 non-null    int64
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB
```

In [4]: `df.isnull().sum()`

```
Out[4]: TransactionID      0
        CustomerID        0
        AccountID         0
        AccountType       0
        TransactionType    0
        Product           0
        Firm              0
        Region            0
        Manager           0
        TransactionDate    0
        TransactionAmount  0
        AccountBalance     0
        RiskScore         0
        CreditRating       0
        TenureMonths      0
        dtype: int64
```

```
In [5]: df.shape
```

```
Out[5]: (800, 15)
```

```
In [6]: df.describe().T
```

```
Out[6]:
```

	count	mean	std	min	max
TransactionID	800.0	98.401250	56.292618	1.000000	50.000
TransactionAmount	800.0	51575.671765	29055.838886	-30721.247890	31692.004
AccountBalance	800.0	72345.657910	34070.038539	-37293.600250	49157.502
RiskScore	800.0	0.473415	0.242858	-0.389354	0.313
CreditRating	800.0	566.586250	151.674261	305.000000	436.500
TenureMonths	800.0	126.582500	68.985320	6.000000	65.000

The dataset contains mixed types (strings, numbers, dates).

Missing values exist in transaction amount and date fields.

Some columns contain commas, symbols, or inconsistent casing.

Numerical distributions show right skewness due to high-value transactio

```
In [7]: df.duplicated().sum()
```

```
Out[7]: 0
```

Data Cleaning Section

Task 1 : Data Cleaning and Formatting

1. Remove/treat any special characters or non-numeric entries from financial fields.

In this step, we clean the dataset by:

1. Formatting inconsistent text
2. Removing special symbols
3. Handling missing values
4. Changing datatypes
5. Standardizing categories

This ensures the dataset is ready for reliable analysis.

```
In [8]: df.columns
```

```
Out[8]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',  
              'TransactionType', 'Product', 'Firm', 'Region', 'Manager',  
              'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',  
              'CreditRating', 'TenureMonths'],  
             dtype='object')
```

```
In [9]: financial_cols = ['TransactionAmount', 'AccountBalance', 'RiskScore',  
                          'CreditRating', 'TenureMonths']  
for col in financial_cols:  
    df[col] = df[col].replace('[^0-9.-]', '', regex=True)  
    df[col] = pd.to_numeric(df[col], errors = 'coerce')
```

```
In [10]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionID          800 non-null    int64
1   CustomerID            800 non-null    object
2   AccountID             800 non-null    object
3   AccountType           800 non-null    object
4   TransactionType        800 non-null    object
5   Product               800 non-null    object
6   Firm                  800 non-null    object
7   Region                800 non-null    object
8   Manager               800 non-null    object
9   TransactionDate        800 non-null    object
10  TransactionAmount      800 non-null    float64
11  AccountBalance         800 non-null    float64
12  RiskScore              800 non-null    float64
13  CreditRating           800 non-null    int64
14  TenureMonths           800 non-null    int64
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB

```

2. Standardize text columns

```

In [11]: text_cols = ['AccountType',
                    'TransactionType', 'Product', 'Firm', 'Region', 'Manager']
for col in text_cols:
    df[col] = df[col].str.strip().str.title()

```

```

In [12]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionID          800 non-null    int64
1   CustomerID            800 non-null    object
2   AccountID             800 non-null    object
3   AccountType           800 non-null    object
4   TransactionType        800 non-null    object
5   Product               800 non-null    object
6   Firm                  800 non-null    object
7   Region                800 non-null    object
8   Manager               800 non-null    object
9   TransactionDate        800 non-null    object
10  TransactionAmount      800 non-null    float64
11  AccountBalance         800 non-null    float64
12  RiskScore              800 non-null    float64
13  CreditRating           800 non-null    int64
14  TenureMonths           800 non-null    int64
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB

```

```

In [13]: df['TransactionAmount'] = df['TransactionAmount'].astype(str).str.replace(r'[$,]', '')
df['TransactionAmount'] = pd.to_numeric(df['TransactionAmount'], errors='coerce')

df['AccountBalance'] = df['AccountBalance'].astype(str).str.replace(r'[$,]', '')
df['AccountBalance'] = pd.to_numeric(df['AccountBalance'], errors='coerce')
print("Invalid Amounts:", df['TransactionAmount'].isnull().sum())

```

Invalid Amounts: 0

Validate and format date columns.

```
In [14]: df.columns
```

```

Out[14]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',
               'TransactionType', 'Product', 'Firm', 'Region', 'Manager',
               'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',
               'CreditRating', 'TenureMonths'],
              dtype='object')

```

The TransactionDate column contains improperly formatted dates. We convert them to a standard datetime format so we can calculate:

Last transaction date

Monthly trends

Account activity levels

Converting Datetime

```
In [15]: df['TransactionDate'] = pd.to_datetime(df['TransactionDate'], errors='coerce', d
```

```
In [16]: df.head(5)
```

```
Out[16]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product
0	33	CUST6549	ACC12334	Credit	Withdrawal	Saving Account
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan

All dates are now in a consistent datetime format.

Invalid or corrupt date entries were converted to NaT and can be cleaned later.

Sorting by date allows accurate activity and recency calculations.

```
In [17]: df['TransactionDate'].isna().sum()
```

```
Out[17]: 0
```

```
In [18]: df.shape
```

```
Out[18]: (800, 15)
```

Feature Engineering

We extract Year and Month from the transaction date to analyze trends over time and group transactions on a monthly basis.

```
In [19]: df['Year'] = df['TransactionDate'].dt.year  
df['Month'] = df['TransactionDate'].dt.month
```

```
In [20]: df.shape
```

```
Out[20]: (800, 17)
```

```
In [21]: df = df.sort_values(by='TransactionDate')
```

Ensure account types and transaction categories are standardized.

We remove extra spaces, inconsistent casing, and formatting issues in AccountType and TransactionType. This ensures accurate grouping and aggregation.

```
In [22]: df['AccountType'] = df['AccountType'].str.strip()
df['TransactionType'] = df['TransactionType'].str.strip()
```

```
In [23]: df['AccountType'] = df['AccountType'].str.title()
df['TransactionType'] = df['TransactionType'].str.title()
```

```
In [24]: df.tail(2).T
```

Out[24]:

	389	33
TransactionID	187	92
CustomerID	CUST3725	CUST5174
AccountID	ACC18177	ACC16664
AccountType	Credit	Credit
TransactionType	Deposit	Transfer
Product	Personal Loan	Mutual Fund
Firm	Firm A	Firm B
Region	Central	South
Manager	Manager 1	Manager 2
TransactionDate	2024-06-22 00:00:00	2024-06-22 00:00:00
TransactionAmount	30610.88038	28562.48672
AccountBalance	30841.45023	101656.4123
RiskScore	0.780858	0.301736
CreditRating	473	355
TenureMonths	37	219
Year	2024	2024
Month	6	6

```
In [25]: df['AccountType'] = df['AccountType'].replace({
    'savings': 'Savings',
    'Curent': 'Current'
})
df['TransactionType'] = df['TransactionType'].replace({
```



```
'Credit':'Credit',  
'Debit':'Debit'  
})
```

```
In [26]: df.columns
```

```
Out[26]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',  
              'TransactionType', 'Product', 'Firm', 'Region', 'Manager',  
              'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',  
              'CreditRating', 'TenureMonths', 'Year', 'Month'],  
              dtype='object')
```

```
In [27]: print(df['AccountType'].unique())  
print(df['TransactionType'].unique())  
print(df['Product'].unique())
```

```
['Loan' 'Savings' 'Credit' 'Current']  
['Withdrawal' 'Transfer' 'Deposit' 'Payment']  
['Credit Card' 'Home Loan' 'Personal Loan' 'Mutual Fund' 'Savings Account']
```

Task 2: Descriptive Transactional Analysis

1. Calculate monthly and yearly summaries of total credits, debits, and net transaction volume.

```
In [28]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 800 entries, 399 to 33
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionID          800 non-null    int64
1   CustomerID            800 non-null    object
2   AccountID             800 non-null    object
3   AccountType           800 non-null    object
4   TransactionType        800 non-null    object
5   Product               800 non-null    object
6   Firm                  800 non-null    object
7   Region                800 non-null    object
8   Manager               800 non-null    object
9   TransactionDate        800 non-null    datetime64[ns]
10  TransactionAmount      800 non-null    float64
11  AccountBalance         800 non-null    float64
12  RiskScore              800 non-null    float64
13  CreditRating           800 non-null    int64
14  TenureMonths           800 non-null    int64
15  Year                   800 non-null    int32
16  Month                  800 non-null    int32
dtypes: datetime64[ns](1), float64(3), int32(2), int64(3), object(8)
memory usage: 106.2+ KB

```

```
In [29]: df['AccountType'].unique()
```

```
Out[29]: array(['Loan', 'Savings', 'Credit', 'Current'], dtype=object)
```

```
In [30]: df['TransactionType'].unique()
```

```
Out[30]: array(['Withdrawal', 'Transfer', 'Deposit', 'Payment'], dtype=object)
```

```
In [31]: # df.drop(columns=['CreditAmount', 'DebitAmount'], inplace=True)
```

```
In [32]: df['TransactionCategory'] = df['TransactionType'].map({
    'Deposit': 'Credit',
    'Withdrawal': 'Debit',
    'Payment': 'Debit',
    'Transfer': 'Debit'
})
df.head(5)
```

```
Out[32]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
399	165	CUST5174	ACC18177	Loan	Withdrawal	Cl
86	121	CUST9248	ACC51971	Savings	Withdrawal	H
606	60	CUST6947	ACC94907	Credit	Transfer	Pers
181	123	CUST7481	ACC21719	Loan	Withdrawal	Mu
377	90	CUST1913	ACC50439	Savings	Withdrawal	Sav Acc

```
In [33]: df['CreditAmount'] = df['TransactionAmount'].where(
        df['TransactionCategory'] == 'Credit', 0
    )

df['DebitAmount'] = df['TransactionAmount'].where(
    df['TransactionCategory'] == 'Debit', 0
)
```

```
In [34]: df['TransactionDate'].dtype
```

```
Out[34]: dtype('<M8[ns]')
```

```
In [35]: df['TransactionDate'] = pd.to_datetime(
        df['TransactionDate'], errors='coerce'
    )
```

```
In [36]: monthly_summary = (
        df.groupby(['Year', 'Month'])
        .agg(
            Total_Credits=('CreditAmount', 'sum'),
            Total_Debits=('DebitAmount', 'sum')
        )
        .reset_index()
    )

monthly_summary['Net_Transaction_Volume'] = (
    monthly_summary['Total_Credits'] -
    monthly_summary['Total_Debits']
)
```

```
In [37]: monthly_summary[['Total_Credits', 'Total_Debits', 'Net_Transaction_Volume']] =
        monthly_summary[['Total_Credits', 'Total_Debits', 'Net_Transaction_Volume']]
```

```
In [38]: monthly_summary['Date'] = pd.to_datetime(
        monthly_summary['Year'].astype(str) + '-' +
```

```
monthly_summary['Month'].astype(str) + '-01'
)
```

```
In [39]: yearly_summary = (
    df.groupby('Year')
      .agg(
        Total_Credits=('CreditAmount', 'sum'),
        Total_Debits=('DebitAmount', 'sum')
      )
      .reset_index()
)

yearly_summary['Net_Transaction_Volume'] = (
    yearly_summary['Total_Credits'] -
    yearly_summary['Total_Debits']
)
```

```
In [40]: yearly_summary[['Total_Credits', 'Total_Debits', 'Net_Transaction_Volume']] =
    yearly_summary[['Total_Credits', 'Total_Debits', 'Net_Transaction_Volume']]
```

```
In [41]: monthly_summary.head()
```

```
Out[41]:
```

	Year	Month	Total_Credits	Total_Debits	Net_Transaction_Volume	Date
0	2023	1	738270.10	2318376.49	-1580106.39	2023-01-01
1	2023	2	648261.00	1433853.41	-785592.41	2023-02-01
2	2023	3	604002.42	2027793.44	-1423791.01	2023-03-01
3	2023	4	439321.69	1338294.29	-898972.60	2023-04-01
4	2023	5	425589.87	2203715.27	-1778125.40	2023-05-01

```
In [42]: yearly_summary
```

```
Out[42]:
```

	Year	Total_Credits	Total_Debits	Net_Transaction_Volume
0	2023	7251368.60	21098815.19	-13847446.59
1	2024	3229767.24	9680586.39	-6450819.15

```
In [43]: yearly_summary['Outflow_Value'] = (
    yearly_summary['Total_Debits'] - yearly_summary['Total_Credits']
)
yearly_summary.round()
```

```
Out[43]:
```

	Year	Total_Credits	Total_Debits	Net_Transaction_Volume	Outflow_Value
0	2023	7251369.0	21098815.0	-13847447.0	13847447.0
1	2024	3229767.0	9680586.0	-6450819.0	6450819.0

```
In [44]: df[['TransactionType', 'TransactionCategory']].drop_duplicates()
```

```
Out[44]:
```

	TransactionType	TransactionCategory
399	Withdrawal	Debit
606	Transfer	Debit
265	Deposit	Credit
689	Payment	Debit

```
In [45]: df['CreditAmount'] = df['TransactionAmount'].where(
    df['TransactionCategory'] == 'Credit', 0
)

df['DebitAmount'] = df['TransactionAmount'].where(
    df['TransactionCategory'] == 'Debit', 0
)
df.head()
```

```
Out[45]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
399	165	CUST5174	ACC18177	Loan	Withdrawal	Ci
86	121	CUST9248	ACC51971	Savings	Withdrawal	H
606	60	CUST6947	ACC94907	Credit	Transfer	Pers
181	123	CUST7481	ACC21719	Loan	Withdrawal	Mu
377	90	CUST1913	ACC50439	Savings	Withdrawal	Sav

```
In [46]: df['TransactionCategory'].value_counts()
```

```
Out[46]: TransactionCategory
Debit      601
Credit     199
Name: count, dtype: int64
```

```
In [47]: df['TransactionType'].value_counts()
```

```
Out[47]: TransactionType
Withdrawal  207
Payment     200
Deposit     199
Transfer    194
Name: count, dtype: int64
```

Transaction types in the dataset (Deposit, Withdrawal, Payment, Transfer) were

mapped to Credit and Debit categories based on financial inflow and outflow behavior. This mapping enabled accurate calculation of total credits, debits, and net transaction volume.

In most months, total debits (₹) are higher than total credits (₹), resulting in a negative net transaction volume.

On a yearly basis, total debits (₹) exceed total credits (₹), leading to an overall net outflow Values As Provided

The net transaction volume remains negative overall, confirming higher cash outflow than inflow across accounts.

```
In [48]: df['TransactionDate'].info()
```

```
<class 'pandas.core.series.Series'>
Index: 800 entries, 399 to 33
Series name: TransactionDate
Non-Null Count  Dtype
-----
800 non-null    datetime64[ns]
dtypes: datetime64[ns](1)
memory usage: 12.5 KB
```

```
In [49]: df.columns
```

```
Out[49]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',
               'TransactionType', 'Product', 'Firm', 'Region', 'Manager',
               'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',
               'CreditRating', 'TenureMonths', 'Year', 'Month', 'TransactionCategory',
               'CreditAmount', 'DebitAmount'],
              dtype='object')
```

Plot trends in total credits vs. debits over time.

```
In [50]: monthly_summary['Date'] = pd.to_datetime(
        monthly_summary['Year'].astype(str) + '-' +
        monthly_summary['Month'].astype(str) + '-01'
    )
```

```
In [51]: monthly_summary = monthly_summary.sort_values('Date')
```


```
In [52]: plt.figure(figsize=(12,6))

plt.plot(
    monthly_summary['Date'],
    monthly_summary['Total_Credits'],
    label='Total Credits',
    linewidth=2
)
```

```
plt.plot(
    monthly_summary['Date'],
    monthly_summary['Total_Debits'],
    label='Total Debits',
    linewidth=2
)

plt.title('Monthly Trend: Total Credits vs Total Debits')
plt.xlabel('Date')
plt.ylabel('Transaction Amount')
plt.legend()
plt.grid(True)

plt.show()
```



The analysis shows that debit transactions consistently exceed credit transactions across most months. This indicates a net cash outflow pattern, suggesting that customers generally withdraw or spend more funds than they deposit during the observed period.

1. Identify top and bottom performing accounts based on net inflow.

Net Inflow Calculation (Credit - Debit)

Net inflow helps us understand whether an account is depositing more or withdrawing more. This is important for identifying:

High value accounts

Accounts showing withdrawal dominance

Accounts at risk of dropping activity

```
In [53]: account_performance = df.groupby('AccountID')[['CreditAmount', 'DebitAmount']].  
account_performance = account_performance.reset_index()
```

```
In [54]: account_performance['Net_Inflow'] = (account_performance['CreditAmount'] - acco
```

```
In [55]: account_performance = account_performance.rename(columns={  
    'CreditAmount': 'Total_Credits',  
    'DebitAmount': 'Total_Debits'  
})
```

```
In [56]: top_accounts = account_performance.sort_values(  
    by='Net_Inflow', ascending=False  
).head(10)  
  
top_accounts
```

```
Out[56]:
```

	AccountID	Total_Credits	Total_Debits	Net_Inflow
92	ACC48501	346856.33960	0.00000	346856.33960
60	ACC33287	390354.42641	201236.66948	189117.75693
168	ACC87006	245497.37832	101488.36862	144009.00970
100	ACC50817	244837.13480	123447.96911	121389.16569
191	ACC99117	167040.52263	45808.22065	121232.30198
112	ACC57516	221384.17912	107219.92732	114164.25180
186	ACC95774	110414.82355	14746.31435	95668.50920
0	ACC10117	142170.20378	57310.76365	84859.44013
86	ACC45951	98505.08125	19454.38267	79050.69858
56	ACC31902	204456.94267	126739.66491	77717.27776

```
In [57]: bottom_accounts = account_performance.sort_values(  
    by='Net_Inflow', ascending=True  
).head(10)  
  
bottom_accounts
```


Out[57]:

	AccountID	Total_Credits	Total_Debits	Net_Inflow
107	ACC53466	18181.67381	476775.484990	-458593.811180
49	ACC29396	39888.00143	442832.014120	-402944.012690
118	ACC60432	39623.16730	423334.682580	-383711.515280
48	ACC29356	27344.62435	408228.871838	-380884.247488
55	ACC31539	25819.70053	376944.670070	-351124.969540
159	ACC82298	0.00000	347376.540660	-347376.540660
165	ACC83581	39304.24834	386292.484270	-346988.235930
185	ACC95164	0.00000	340602.161780	-340602.161780
13	ACC16241	103643.14455	435968.998300	-332325.853750
99	ACC50439	0.00000	332311.198690	-332311.198690

Net Inflow Insights

Accounts with positive net inflow are likely depositing more money and show healthier engagement.

Accounts with highly negative net inflow may be withdrawing or becoming inactive.

Top 25% inflow accounts can be considered high-value customers.

There is a clear contrast between top and bottom accounts, with top performers maintaining surplus balances while bottom performers experience sustained cash depletion.

In [58]: `df.head(2)`

Out[58]:

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
399	165	CUST5174	ACC18177	Loan	Withdrawal	C
86	121	CUST9248	ACC51971	Savings	Withdrawal	H

Identify and flag accounts as dormant or inactive if there is a gap of two months or more between consecutive transactions.

In [59]: `df = df.sort_values(['AccountID', 'TransactionDate'])`

In [60]: `df['Prev_Transaction_Date'] = df.groupby('AccountID')['TransactionDate'].shift(1)`
`df['Gap_Days'] = (`

```
df['TransactionDate'] - df['Prev_Transaction_Date']  
) .dt.days
```

```
In [61]: df.columns
```

```
Out[61]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',  
              'TransactionType', 'Product', 'Firm', 'Region', 'Manager',  
              'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',  
              'CreditRating', 'TenureMonths', 'Year', 'Month', 'TransactionCategor  
y',  
              'CreditAmount', 'DebitAmount', 'Prev_Transaction_Date', 'Gap_Days'],  
              dtype='object')
```

```
In [62]: # df['Gap_Months'] = df['Gap_Days'] / 30
```

```
In [63]: # dormant_accounts = (  
#         df.groupby('AccountID')['Gap_Months']  
#         .max()  
#         .reset_index()  
# )  
  
# dormant_accounts['Account_Status'] = dormant_accounts['Gap_Months'].apply(  
#     lambda x: 'Dormant' if x >= 2 else 'Active'  
# )
```

```
In [64]: # df = df.merge(  
#         dormant_accounts[['AccountID', 'Account_Status']],  
#         on='AccountID',  
#         how='left'  
# )
```

```
In [65]: # df['Account_Status'].value_counts()
```

```
In [ ]: ### As we can see There are 53 Active accounts and 747 Dormant Account
```

Defining Active, Inactive, and Dormant Accounts

To measure customer engagement, we create a new column AccountStatus using the following rules:

✓ Activity Criteria:

Active → Last transaction within 30 days

Inactive → Last transaction within 30–90 days

Dormant → No transaction for 90+ days

These definitions help the bank:

Identify churn risk

Target customers with offers

Prioritize high-engagement users

```
In [153... import datetime as dt

# 2. Recalculate last transaction date per account
last_txn = df.groupby('AccountID')['TransactionDate'].max().reset_index()
last_txn.rename(columns={'TransactionDate': 'Last_Transaction'}, inplace=True)

# 3. Merge
df = df.merge(last_txn, on='AccountID', how='left')

# 4. Calculate days since last transaction using today's date
today = df['TransactionDate'].max()
df['Days_Since_Last_Txn'] = (today - df['Last_Transaction']).dt.days

# 5. Define activity
def activity_status(days):
    if pd.isna(days):
        return 'Dormant'
    if days <= 30:
        return 'Active'
    elif days <= 90:
        return 'Inactive'
    else:
        return 'Dormant'

df['AccountStatus'] = df['Days_Since_Last_Txn'].apply(activity_status)
```

```
In [154... df['AccountStatus'].value_counts()
```

```
Out[154... AccountStatus
Dormant      420
Active       235
Inactive     145
Name: count, dtype: int64
```

Account Activity Distribution

After applying the activity rules based on transaction recency:

Active → Last transaction within 30 days

Inactive → Last transaction within 31–90 days

Dormant → No transactions in the last 90+ days

The distribution of customer activity levels is:

Dormant Accounts: 420

Active Accounts: 235

Inactive Accounts: 145

Insights

The largest segment (420 accounts) falls under Dormant, indicating long periods without transactions. These accounts may require customer re-engagement strategies.

Active accounts (235) show regular and recent engagement, indicating healthy usage behavior.

Inactive accounts (145) represent medium-risk customers who have not transacted recently but can still be reactivated through targeted campaigns.

Business Implications

Dormant customers may be at risk of churn; targeted communication or promotional offers may help reactivate them.

Active customers could benefit from premium financial products or personalized recommendations.

Inactive customers should be monitored closely and nudged with reminders or early interventions to prevent inactivity from turning into dormancy.

Grouping accounts by activity levels: High, Medium, Low based on transaction frequency on your analysis and rubrics. Do not forget to mention the rubric in the headings.

Account Activity Level Segmentation (Rubric: High $\geq$ 10 | Medium 5–9 | Low $<$ 5 transactions)

In [156...

df

Out[156...

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
0	73	CUST9209	ACC10117	Savings	Transfer	Cl
1	34	CUST3109	ACC10117	Savings	Deposit	H
2	106	CUST8155	ACC10117	Credit	Deposit	Mu
3	169	CUST3810	ACC10117	Loan	Deposit	H
4	158	CUST3041	ACC10996	Savings	Withdrawal	Cl
...	...	...	...	...	...	
795	45	CUST7887	ACC99409	Current	Withdrawal	H
796	135	CUST3015	ACC99549	Savings	Transfer	Cl
797	195	CUST8416	ACC99549	Loan	Payment	Cl
798	25	CUST4258	ACC99549	Loan	Transfer	Mu
799	92	CUST5385	ACC99549	Loan	Deposit	Mu

800 rows × 36 columns

```
In [69]: account_activity = (
    df.groupby('AccountID')
      .size()
      .reset_index(name='Transaction_Count')
    )
```

```
In [70]: def activity_level(count):
    if count >= 10:
        return 'High'
    elif count >= 5:
        return 'Medium'
    else:
        return 'Low'
```

```
In [71]: account_activity['Activity_Level'] = (
    account_activity['Transaction_Count']
      .apply(activity_level)
    )
```

```
In [72]: account_activity['Activity_Level'].value_counts()
```

```
Out[72]: Activity_Level
Low      121
Medium   70
High      3
Name: count, dtype: int64
```

```
In [73]: df = df.merge(
    account_activity[['AccountID', 'Activity_Level']],
    on='AccountID',
    how='left'
)
```

```
In [74]: df.columns
```

```
Out[74]: Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',
    'TransactionType', 'Product', 'Firm', 'Region', 'Manager',
    'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',
    'CreditRating', 'TenureMonths', 'Year', 'Month', 'TransactionCategor
y',
    'CreditAmount', 'DebitAmount', 'Prev_Transaction_Date', 'Gap_Days',
    'Last_Transaction', 'Days_Since_Last_Txn', 'AccountStatus',
    'Activity_Level'],
    dtype='object')
```

```
In [75]: df[['AccountID', 'Activity_Level']].head()
```

```
Out[75]:
```

	AccountID	Activity_Level
0	ACC10117	Low
1	ACC10117	Low
2	ACC10117	Low
3	ACC10117	Low
4	ACC10996	Medium

```
In [158]: # After creating freq_balance
transaction_threshold_high = freq_balance['Transaction_Count'].quantile(0.75)
transaction_threshold_low = freq_balance['Transaction_Count'].quantile(0.25)

balance_threshold_low = freq_balance['Avg_Balance'].quantile(0.25)

# Create frequency category
def freq_category(n):
    if n >= transaction_threshold_high:
        return 'High Frequency'
    elif n <= transaction_threshold_low:
        return 'Low Frequency'
    else:
        return 'Medium Frequency'

freq_balance['Freq_Category'] = freq_balance['Transaction_Count'].apply(freq_c
```

```
# Create balance category
freq_balance['Balance_Category'] = freq_balance['Avg_Balance'].apply(
    lambda x: 'Low Balance' if x <= balance_threshold_low else 'Normal Balance'
)
```

```
In [161]: df = df.merge(freq_balance[['AccountID', 'Freq_Category', 'Balance_Category']],
                        on='AccountID', how='left')
```

```
In [162]: df
```

```
Out[162]:
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
0	73	CUST9209	ACC10117	Savings	Transfer	Cl
1	34	CUST3109	ACC10117	Savings	Deposit	H
2	106	CUST8155	ACC10117	Credit	Deposit	Mu
3	169	CUST3810	ACC10117	Loan	Deposit	H
4	158	CUST3041	ACC10996	Savings	Withdrawal	Cl
...	...	...	...	...	...	...
795	45	CUST7887	ACC99409	Current	Withdrawal	H
796	135	CUST3015	ACC99549	Savings	Transfer	Cl
797	195	CUST8416	ACC99549	Loan	Payment	Cl
798	25	CUST4258	ACC99549	Loan	Transfer	Mu
799	92	CUST5385	ACC99549	Loan	Deposit	Mu

800 rows × 40 columns

Segment customers by average balance and transaction volume.

```
In [76]: customer_summary = (
    df.groupby('AccountID')
    .agg(
        Avg_Balance=('AccountBalance', 'mean'),
        Transaction_Volume=('TransactionAmount', 'sum')
    )
)
```

```
)
    .reset_index()
)
```

```
In [77]: balance_threshold = customer_summary['Avg_Balance'].median()
volume_threshold = customer_summary['Transaction_Volume'].median()
```

```
In [78]: def segment_customer(row):
    if row['Avg_Balance'] >= balance_threshold and row['Transaction_Volume'] >
        return 'High Value - Active'
    elif row['Avg_Balance'] >= balance_threshold and row['Transaction_Volume']
        return 'High Balance - Low Usage'
    elif row['Avg_Balance'] < balance_threshold and row['Transaction_Volume']
        return 'Low Balance - High Usage'
    else:
        return 'Low Value - Inactive'
```

```
In [79]: customer_summary['Customer_Segment'] = customer_summary.apply(
    segment_customer, axis=1
)
```

```
In [80]: df = df.merge(
    customer_summary[['AccountID', 'Customer_Segment']],
    on='AccountID',
    how='left'
)
```

```
In [81]: customer_summary['Customer_Segment'].value_counts()
```

```
Out[81]: Customer_Segment
Low Value - Inactive      49
High Value - Active      49
Low Balance - High Usage  48
High Balance - Low Usage  48
Name: count, dtype: int64
```

Create profiles for:

1. High-net inflow accounts
2. High-frequency low-balance accounts
3. Accounts with negative or near-zero balances

```
In [82]: account_inflow = (
    df.groupby('AccountID').agg(
        Total_Credits = ('CreditAmount', 'sum'),
        Total_Debits = ('DebitAmount', 'sum')
    ).reset_index()
)
account_inflow['Net_Inflow'] = (
    account_inflow['Total_Credits'] - account_inflow['Total_Debits']
)
```



```
In [83]: high_inflow_threshold = account_inflow['Net_Inflow'].quantile(0.75)
high_net_inflow_accounts = account_inflow[
account_inflow['Net_Inflow'] > high_inflow_threshold
]
```

```
In [84]: freq_balance = (
df.groupby('AccountID')
    .agg(
        Transaction_Count=('TransactionAmount', 'count'),
        Avg_Balance=('AccountBalance', 'mean')
    )
    .reset_index()
)
```

```
In [85]: freq_threshold = freq_balance['Transaction_Count'].quantile(0.75)
balance_threshold = freq_balance['Avg_Balance'].quantile(0.25)
```

```
In [86]: high_freq_low_balance_accounts = freq_balance[
    (freq_balance['Transaction_Count'] >= freq_threshold) &
    (freq_balance['Avg_Balance'] <= balance_threshold)
]
```

```
In [87]: negative_near_zero_accounts = freq_balance[
    freq_balance['Avg_Balance'] <= 5000
]
```

```
In [88]: high_net_inflow_accounts.head(5)
```

```
Out[88]:
```

	AccountID	Total_Credits	Total_Debits	Net_Inflow
0	ACC10117	142170.20378	57310.76365	84859.44013
2	ACC11062	0.00000	4014.26490	-4014.26490
7	ACC12334	149116.79761	161110.49208	-11993.69447
10	ACC15359	122856.48300	120413.19788	2443.28512
17	ACC18140	45680.91092	38542.74767	7138.16325

```
In [89]: print("High-Net Inflow Accounts:", high_net_inflow_accounts.shape[0])
print("High-Frequency Low-Balance Accounts:", high_freq_low_balance_accounts.s
print("Negative / Near-Zero Balance Accounts:", negative_near_zero_accounts.sh
```

High-Net Inflow Accounts: 49

High-Frequency Low-Balance Accounts: 15

Negative / Near-Zero Balance Accounts: 1

Customer profiles were created to capture key financial behavior patterns. High-net inflow accounts were identified using upper-quartile net inflow thresholds. High-frequency low-balance accounts were flagged using transaction frequency and balance percentiles. Accounts with negative or near-zero balances were highlighted to assess financial stress and overdraft risk. These profiles support targeted engagement and proactive risk monitoring.

Track accounts with frequent large withdrawals or overdrafts.

```
In [90]: withdrawals = df[df['TransactionType'] == 'Withdrawal'].copy()
```

```
In [91]: large_withdrawal_threshold = withdrawals['TransactionAmount'].quantile(0.75)
print("Highest Large Withdrawals or Overdrafts : ₹",large_withdrawal_threshold)
```

Highest Large Withdrawals or Overdrafts : ₹ 73511.85

```
In [92]: withdrawals['Large_Withdrawal'] = (
    withdrawals['TransactionAmount'] >= large_withdrawal_threshold
)
```

```
In [93]: large_withdrawal_counts = (
    withdrawals[withdrawals['Large_Withdrawal']]
    .groupby('AccountID')
    .size()
    .reset_index(name='Large_Withdrawal_Count')
)
```

```
In [94]: frequent_large_withdrawal_accounts = large_withdrawal_counts[
    large_withdrawal_counts['Large_Withdrawal_Count'] >= 3
]
```

frequent_large_withdrawal_accounts

```
Out[94]:
```

	AccountID	Large_Withdrawal_Count
--	-----------	------------------------

```
In [95]: df['Overdraft'] = df['AccountBalance'] < 0
```

```
In [96]: overdraft_counts = (
    df[df['Overdraft']]
    .groupby('AccountID')
    .size()
    .reset_index(name='Overdraft_Count')
)
```

```
In [97]: frequent_overdraft_accounts = overdraft_counts[
    overdraft_counts['Overdraft_Count'] >= 2
]
```

frequent_overdraft_accounts

```
Out[97]:
```

	AccountID	Overdraft_Count
--	-----------	-----------------

5	ACC28292	2
---	----------	---

```
In [98]: risky_accounts = pd.merge(
    frequent_large_withdrawal_accounts,
```

```

    frequent_overdraft_accounts,
    on='AccountID',
    how='outer'
).fillna(0)

risky_accounts

```

Out[98]:

	AccountID	Large-Withdrawal_Count	Overdraft_Count
0	ACC28292	0.0	2

```

In [99]: df['Frequent_Large-Withdrawal'] = df['AccountID'].isin(
    frequent_large_withdrawal_accounts['AccountID']
)

df['Frequent_Overdraft'] = df['AccountID'].isin(
    frequent_overdraft_accounts['AccountID']
)

```

```

In [100]: df[['AccountID', 'Frequent_Large-Withdrawal', 'Frequent_Overdraft']].drop_duplicates()

```

Out[100]:

	AccountID	Frequent_Large-Withdrawal	Frequent_Overdraft
0	ACC10117	False	False
4	ACC10996	False	False
9	ACC11062	False	False
11	ACC11188	False	False
16	ACC11285	False	False

```

In [101]: print("Accounts with frequent large withdrawals:",
    frequent_large_withdrawal_accounts.shape[0])

print("Accounts with frequent overdrafts:",
    frequent_overdraft_accounts.shape[0])

```

Accounts with frequent large withdrawals: 0
Accounts with frequent overdrafts: 1

Calculate balance volatility using standard deviation or coefficient of variation.

```

In [102]: balance_volatility = (
    df.groupby('AccountID')
    .agg(
        Avg_Balance=('AccountBalance', 'mean'),
        Balance_STD=('AccountBalance', 'std')
    )
    .reset_index()
)

```

```
In [103... balance_volatility['Balance_STD'] = balance_volatility['Balance_STD'].fillna(0)
```

```
In [104... balance_volatility['Balance_CV'] = (  
    balance_volatility['Balance_STD'] / balance_volatility['Avg_Balance'].abs(  
    )  
)
```

```
In [105... balance_volatility.replace([float('inf'), -float('inf')], 0, inplace=True)  
balance_volatility['Balance_CV'] = balance_volatility['Balance_CV'].fillna(0)
```

```
In [106... cv_threshold = balance_volatility['Balance_CV'].quantile(0.75)  
  
balance_volatility['High_Balance_Volatility'] = (  
    balance_volatility['Balance_CV'] >= cv_threshold  
)
```

```
In [107... high_volatility_accounts = balance_volatility[  
    balance_volatility['High_Balance_Volatility']  
    ]  
  
high_volatility_accounts.head(7)
```

```
Out[107... 
```

	AccountID	Avg_Balance	Balance_STD	Balance_CV	High_Balance_Volatility
4	ACC11285	97401.348560	55922.732441	0.574147	True
5	ACC11837	84852.733695	60694.391957	0.715291	True
8	ACC13357	69179.806513	40614.664902	0.587088	True
13	ACC16241	73521.710375	45312.204045	0.616311	True
17	ACC18140	39960.076965	31325.244170	0.783914	True
25	ACC21878	67726.446770	70517.525297	1.041211	True
29	ACC23736	60801.533102	48785.760924	0.802377	True

```
In [108... df = df.merge(  
    balance_volatility[['AccountID', 'Balance_STD', 'Balance_CV', 'High_Balance_Volatility']],  
    on='AccountID',  
    how='left'  
)
```

```
In [109... df[['AccountID', 'Balance_STD', 'Balance_CV', 'High_Balance_Volatility']].drop
```

Out[109...	AccountID	Balance_STD	Balance_CV	High_Balance_Volatility
0	ACC10117	25886.972758	0.369249	False
4	ACC10996	9434.002316	0.216535	False
9	ACC11062	3208.737888	0.084137	False
11	ACC11188	35494.660810	0.509599	False
16	ACC11285	55922.732441	0.574147	True

Use IQR or z-score methods to detect anomalies

In [110... *# Using IQR method*

```
Q1 = df['TransactionAmount'].quantile(0.25)
Q3 = df['TransactionAmount'].quantile(0.75)

IQR = Q3 - Q1
```

In [111... lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

In [112... lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

In [113... df['IQR_Anomaly'] = (
 (df['TransactionAmount'] < lower_bound) |
 (df['TransactionAmount'] > upper_bound)
)

In [114... iqr_anomalies = df[df['IQR_Anomaly'] == True]
iqr_anomalies.head(10)

Out[114...	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Proc
103	14	CUST3015	ACC21719	Loan	Deposit	Sav Acc

1 rows × 34 columns

In [115... iqr_anomaly_counts = (
 iqr_anomalies.groupby('AccountID')
 .size()
 .reset_index(name='IQR_Anomaly_Count')
)
iqr_anomaly_counts.head()

```
Out[115... AccountID IQR_Anomaly_Count
```

	AccountID	IQR_Anomaly_Count
0	ACC21719	1

Highlight customers with irregular or suspicious transaction behavior.

```
In [116... accounts = df[['AccountID']].drop_duplicates()
```

```
In [117... accounts['Frequent_Large-Withdrawal'] = accounts['AccountID'].isin(
    frequent_large_withdrawal_accounts['AccountID']
)
```

```
In [118... accounts['Frequent_Overdraft'] = accounts['AccountID'].isin(
    frequent_overdraft_accounts['AccountID']
)
```

```
In [119... accounts = accounts.merge(
    balance_volatility[['AccountID', 'High_Balance_Volatility']],
    on='AccountID',
    how='left'
)

accounts['High_Balance_Volatility'] = (
    accounts['High_Balance_Volatility'].fillna(False)
)
```

```
In [120... accounts = accounts.merge(
    iqr_anomaly_counts,
    on='AccountID',
    how='left'
)

accounts['IQR_Anomaly_Count'] = accounts['IQR_Anomaly_Count'].fillna(0)

accounts['High_Anomaly'] = accounts['IQR_Anomaly_Count'] >= 2
```

```
In [121... accounts['Risk_Score'] = (
    accounts['Frequent_Large-Withdrawal'].astype(int) +
    accounts['Frequent_Overdraft'].astype(int) +
    accounts['High_Balance_Volatility'].astype(int) +
    accounts['High_Anomaly'].astype(int)
)
```

```
In [122... accounts['Suspicious_Account'] = accounts['Risk_Score'] >= 2
```


```
In [123... suspicious_accounts = accounts[accounts['Suspicious_Account']]
suspicious_accounts.head(5)
```

Out[123...	AccountID	Frequent_Large-Withdrawal	Frequent_Overdraft	High_Balance_V
42	ACC28292	False	True	

```
In [124... suspicious_accounts['Risk_Score'].value_counts()
```

```
Out[124... Risk_Score
2      1
Name: count, dtype: int64
```


```
In [125... plt.figure(figsize=(10,5))
sns.boxplot(x=df['TransactionAmount'])
plt.title('IQR-based Outlier Detection for Transaction Amount')
plt.xlabel('Transaction Amount')
plt.show()
```



```
In [126... plt.figure(figsize=(10,5))

plt.scatter(
    df.index,
    df['TransactionAmount'],
    c=df['IQR_Anomaly'].map({True: 'red', False: 'blue'}),
    alpha=0.6
)

plt.title('IQR Anomaly Detection: Normal vs Anomalous Transactions')
plt.xlabel('Transaction Index')
plt.ylabel('Transaction Amount')
plt.show()
```



IQR-Based Anomaly Detection Visualization


The Interquartile Range (IQR) method was used to identify anomalous transaction amounts.

Visualizations were created to highlight extreme values and compare normal and anomalous transactions.

Conduct extensive exploratory data analysis with attractive visualizations for your findings


To understand how transaction values are spread and identify skewness or extreme values.

```
In [127... plt.figure(figsize=(10,5))
sns.histplot(df['TransactionAmount'],bins=50, kde=True)
plt.title('Distribution of Transaction Amounts')
plt.xlabel('Transaction Amount')
plt.ylabel('Frequency')
plt.show()
```

To understand which transaction types dominate customer activity.

```
In [128... plt.figure(figsize=(8,5))
sns.countplot(x='TransactionType',data=df)
plt.title('Transaction Type Distribution')
plt.xlabel('Transaction Type')
plt.ylabel('Count')
plt.show()
```




To visualize inflow vs outflow trends over time.

```
In [129... plt.figure(figsize=(12,6))

plt.plot(
    monthly_summary['Date'],
    monthly_summary['Total_Credits'],
    label='Total Credits'
)

plt.plot(
    monthly_summary['Date'],
    monthly_summary['Total_Debits'],
    label='Total Debits'
)

plt.legend()
plt.title('Monthly Credit vs Debit Trend')
plt.xlabel('Date')
plt.ylabel('Amount')
plt.show()
```




In [130... `df.columns`

Out[130... Index(['TransactionID', 'CustomerID', 'AccountID', 'AccountType',
 'TransactionType', 'Product', 'Firm', 'Region', 'Manager',
 'TransactionDate', 'TransactionAmount', 'AccountBalance', 'RiskScore',
 'CreditRating', 'TenureMonths', 'Year', 'Month', 'TransactionCategor
 y',
 'CreditAmount', 'DebitAmount', 'Prev_Transaction_Date', 'Gap_Days',
 'Last_Transaction', 'Days_Since_Last_Txn', 'AccountStatus',
 'Activity_Level', 'Customer_Segment', 'Overdraft',
 'Frequent_Large-Withdrawal', 'Frequent_Overdraft', 'Balance_STD',
 'Balance_CV', 'High_Balance_Volatility', 'IQR_Anomaly'],
 dtype='object')


Rubric: High ≥ 10 | Medium 5-9 | Low < 5 transactions

In [131... `plt.figure(figsize=(8,5))`
`sns.countplot(x='Activity_Level', data=df)`
`plt.title('Account Activity Level Distribution')`
`plt.xlabel('Activity Level')`
`plt.ylabel('Number of Transactions')`
`plt.show()`




Average Balance vs Transaction Volume

```
In [132... plt.figure(figsize=(10,6))
sns.scatterplot(
    x='Transaction_Volume',
    y='Avg_Balance',
    hue='Customer_Segment',
    data=customer_summary
)
plt.title('Customer Segmentation by Balance and Transaction Volume')
plt.xlabel('Transaction Volume')
plt.ylabel('Average Balance')
plt.show()
```




Balance Volatility Distribution

```
In [133... plt.figure(figsize=(10,5))
sns.histplot(balance_volatility['Balance_CV'], bins=40, kde=True)
plt.title('Distribution of Balance Volatility (CV)')
plt.xlabel('Coefficient of Variation')
plt.ylabel('Frequency')
plt.show()
```



Anomalous Transactions Visualization


```
In [134... plt.figure(figsize=(10,5))
sns.boxplot(x=df['TransactionAmount'])
plt.title('Transaction Amount Outliers (IQR)')
plt.show()
```



Risk Indicator Contribution

```
In [135... risk_counts = suspicious_accounts[
    ['Frequent_Large-Withdrawal', 'Frequent_Overdraft',
     'High_Balance_Volatility', 'High_Anomaly']
].sum()

plt.figure(figsize=(8,5))
risk_counts.plot(kind='bar')
plt.title('Risk Indicator Contribution')
plt.ylabel('Number of Accounts')
plt.show()
```



Hypothesis Testing

```
In [136... account_stats = (  
    df.groupby('AccountID')  
    .agg(  
        Avg_Balance=('AccountBalance', 'mean'),  
        Transaction_Volume=('TransactionAmount', 'sum')  
    )  
    .reset_index()  
)
```

```
In [137... volume_threshold = account_stats['Transaction_Volume'].median()  
  
high_volume_accounts = account_stats[  
    account_stats['Transaction_Volume'] > volume_threshold  
]['Avg_Balance']  
  
low_volume_accounts = account_stats[
```

```
account_stats['Transaction_Volume'] <= volume_threshold  
]['Avg_Balance']
```

```
In [138... ## Performing Hypothesis Testing  
from scipy.stats import ttest_ind
```

```
In [139... t_stat, p_value = ttest_ind(  
    high_volume_accounts,  
    low_volume_accounts,  
    equal_var=False # Welch's t-test (safer)  
)  
t_stat, p_value
```

```
Out[139... (-0.08733149845245664, 0.9305072815617675)
```

```
In [140... alpha = 0.05  
  
if p_value < alpha:  
    print("Reject the null hypothesis: High-volume accounts have significantly  
else:  
    print("Fail to reject the null hypothesis: No significant difference in av
```


Fail to reject the null hypothesis: No significant difference in average balances.

```
In [141... print("High-volume average balance: ₹", high_volume_accounts.mean().round(2))  
print("Low-volume average balance: ₹", low_volume_accounts.mean().round(2))
```

High-volume average balance: ₹ 71871.43
Low-volume average balance: ₹ 72118.95

Visual Comparison

```
In [142... plot_data = account_stats.copy()  
plot_data['Volume_Group'] = plot_data['Transaction_Volume'].apply(  
    lambda x: 'High Volume' if x > volume_threshold else 'Low Volume'  
)  
  
plt.figure(figsize=(8,5))  
sns.boxplot(x='Volume_Group', y='Avg_Balance', data=plot_data)  
plt.title('Average Balance: High vs Low Transaction Volume Accounts')  
plt.show()
```

Conduct hypothesis testing based on segmentation.

Hypothesis testing was conducted to validate differences observed across customer segments. Accounts were segmented based on transaction volume, and an independent t-test was applied to compare average balances between high-volume and low-volume groups. This approach ensured that observed differences were statistically significant and not due to random variation.

In []: